







Thin: paired patient-cluster 95% CI; thick: ± 1 training-seed SD (n=3). Six fixed noise realizations; no new significance tests.