


Training and clean-validation trajectories | All registered epochs



Thin lines show individual training seeds and thick lines their arithmetic mean; no CI. Training losses involve strategy-specific augmented inputs and are not the checkpoint-selection endpoint. Checkpoints are selected by clean-validation macro AUROC, after all fixed epochs; no test-driven epoch or strategy selection.