







0 dB structure effects | both architectures are co-primary

Training-seed estimates S_00 S_01 S_02 S_03 S_04



Thin: paired patient 95% CI, conditional on fixed checkpoints/noises/modes. Thick: +/-1 training-seed SD (3 seeds); these are separate, not combined.
Positive AUROC/F1: first condition better; positive ECE: first condition worse. E = Standard; S = equal five-mode mean; I = independent RMS.
Deterministic row offsets: CI -0.065 / SD +0.065; seed ticks -0.23; individual mode points +0.23 (small fixed spacing).
Modes are fixed descriptive controls, not extra training seeds. E-I has no mode dependence and is shown once. Grey row: E-S.